

Ch. 7 Data Cleaning and preparation

pg. 191 - 251

```
In [1]: import pandas as pd
import numpy as np
```

In general:

<https://realpython.com/python-map-function/>

A helpful link for this entire chapter, that explains what Mapping, Filtering and Reducing are.

7.1 Handling Missing Data

The standard in pandas is to use NaN for floating pt values.

```
In [7]: string_data = pd.Series(['aardvark', 'artichoke', np.nan, 'avocado'])
string_data
```

```
Out[7]: 0    aardvark
1    artichoke
2         NaN
3     avocado
dtype: object
```

```
In [8]: string_data.isnull()
```

```
Out[8]: 0    False
1    False
2     True
3    False
dtype: bool
```

pandas use of NA/not available is adopted from R.

There is also the built-in Python 'None' value which is also treated like an NA object array.

```
In [10]: string_data[0] = None # We are replacing the FIRST value in string_data with
        an NaN equivalent.
        string_data.isnull() # We ask how many values are NaN now.
```

```
Out[10]: 0    True
        1    False
        2    True
        3    False
        dtype: bool
```

See Tabel 7-1 on pg 192 for a table for some other NA methods. E.g. - fillna, fill the missing data with some value.

Filtering out missing data

Can do this by using pandas.isnull and boolean indexing, OR by using 'dropna'

```
In [11]: # FOR A Series:

        from numpy import nan as NA

        data = pd.Series([1, NA, 3.5, NA, 7]) #Create the series
        data
```

```
Out[11]: 0    1.0
        1    NaN
        2    3.5
        3    NaN
        4    7.0
        dtype: float64
```

```
In [12]: # Now drop all NA values from the series
        data.dropna()
```

```
Out[12]: 0    1.0
        2    3.5
        4    7.0
        dtype: float64
```

```
In [13]: # You can do the above .dropna() OR you can do:
        data[data.notnull()]
```

```
Out[13]: 0    1.0
        2    3.5
        4    7.0
        dtype: float64
```

FOR DATAFRAMES: You may want to drop rows or columns. Using `.dropna` by default drops any row containing a missing value.

```
In [15]: data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA],
                             [NA, NA, NA], [NA, 6.5, 3.]])
data
```

Out[15]:

| | 0 | 1 | 2 |
|---|-----|-----|-----|
| 0 | 1.0 | 6.5 | 3.0 |
| 1 | 1.0 | NaN | NaN |
| 2 | NaN | NaN | NaN |
| 3 | NaN | 6.5 | 3.0 |

```
In [17]: cleaned = data.dropna() # deleting all ROWS with NA values
cleaned
```

Out[17]:

| | 0 | 1 | 2 |
|---|-----|-----|-----|
| 0 | 1.0 | 6.5 | 3.0 |

However, using `how='all'` will only drop rows where ALL its values are NA

```
In [18]: data.dropna(how='all') # So only row 2 was dropped since it had ALL NaN values.
```

Out[18]:

| | 0 | 1 | 2 |
|---|-----|-----|-----|
| 0 | 1.0 | 6.5 | 3.0 |
| 1 | 1.0 | NaN | NaN |
| 3 | NaN | 6.5 | 3.0 |

```
In [19]: # To drop columns the same way as above, use axis = 1
data[4] = NA # This says, add a column '4', and make all values NA
data
```

Out[19]:

| | 0 | 1 | 2 | 4 |
|---|-----|-----|-----|-----|
| 0 | 1.0 | 6.5 | 3.0 | NaN |
| 1 | 1.0 | NaN | NaN | NaN |
| 2 | NaN | NaN | NaN | NaN |
| 3 | NaN | 6.5 | 3.0 | NaN |

```
In [21]: # Next, drop columns there ALL values are NaN
data.dropna(axis = 1, how = 'all')

# MUST put in the axis =1, otherwise only rows with all NAs will be dropped, not columns.
```

Out[21]:

| | 0 | 1 | 2 |
|---|-----|-----|-----|
| 0 | 1.0 | 6.5 | 3.0 |
| 1 | 1.0 | NaN | NaN |
| 2 | NaN | NaN | NaN |
| 3 | NaN | 6.5 | 3.0 |

It is also common to filter rows of a DataFrame from time series data. Suppose you want to keep only rows containing a certain no. of observation. You can indicate this by use the **thresh** argument.

```
In [26]: df = pd.DataFrame(np.random.rand(7, 3)) # Create the dataframe.
df.iloc[:4, 1] = NA # iloc is an index selector, saying replace up the first
4 rows, in column '1' with NaN
df.iloc[:2, 2] = NA # replace all values up to the 3rd row of the column '2' w
ith NA
df
```

Out[26]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | NaN | NaN |
| 1 | 0.547567 | NaN | NaN |
| 2 | 0.125879 | NaN | 0.413911 |
| 3 | 0.397435 | NaN | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

```
In [28]: df.dropna() # Get rid of anything row/column with a NA value.
```

Out[28]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

```
In [29]: # Instead of the above, can do the following if tou want keep certain rows?
df.dropna(thresh = 2)

#The above reads, drop all NA values that are in col '2'
# Notice that in turn this eliminated two rows as well, cuz otherwise the dataf
rame would be imbalanced.
```

Out[29]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 2 | 0.125879 | NaN | 0.413911 |
| 3 | 0.397435 | NaN | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

Filtering *in* missing data

You may want to fill in the 'holes' in the data. The most common method is `fillna`, with a constant to replace the missing values with that given constant.

In [30]:

```
df
```

Out[30]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | NaN | NaN |
| 1 | 0.547567 | NaN | NaN |
| 2 | 0.125879 | NaN | 0.413911 |
| 3 | 0.397435 | NaN | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

In [31]: `df.fillna(0) # Fill in all the NAs in df with a 0`

Out[31]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | 0.000000 | 0.000000 |
| 1 | 0.547567 | 0.000000 | 0.000000 |
| 2 | 0.125879 | 0.000000 | 0.413911 |
| 3 | 0.397435 | 0.000000 | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

You can also call `fillna` with a dict, if you want to have a different fill value for each column!

In [32]: `df # What it looks like before`

Out[32]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | NaN | NaN |
| 1 | 0.547567 | NaN | NaN |
| 2 | 0.125879 | NaN | 0.413911 |
| 3 | 0.397435 | NaN | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

In [33]: `# What it looks like after
df.fillna({1: 0.5, 2:0}) # For col '1', replace all NA with 0.5, and for col
'2' replace all NAs with 0`

Out[33]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | 0.500000 | 0.000000 |
| 1 | 0.547567 | 0.500000 | 0.000000 |
| 2 | 0.125879 | 0.500000 | 0.413911 |
| 3 | 0.397435 | 0.500000 | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

[!] WTF does this mean?? pg. 195

`fillna` returns a `_new object_`, but you can modify the existing object in place...
(??)

```
In [39]: _ = df.fillna(0, inplace = True) # It's the same output if False?
df
```

Out[39]:

| | 0 | 1 | 2 |
|---|----------|----------|----------|
| 0 | 0.325750 | 0.000000 | 0.000000 |
| 1 | 0.547567 | 0.000000 | 0.000000 |
| 2 | 0.125879 | 0.000000 | 0.413911 |
| 3 | 0.397435 | 0.000000 | 0.525204 |
| 4 | 0.374344 | 0.207910 | 0.446015 |
| 5 | 0.834835 | 0.098177 | 0.026528 |
| 6 | 0.261710 | 0.499677 | 0.554304 |

```
In [40]: # You can use the same redindexing methods with fillna:
df = pd.DataFrame(np.random.randn(6, 3)) #creating a new data from, 6x3 with r
anom no.
df
```

Out[40]:

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 0 | -0.639978 | -0.799902 | -0.070594 |
| 1 | 0.725652 | 0.243492 | 0.511909 |
| 2 | -0.386548 | 0.867206 | 0.272159 |
| 3 | -0.475659 | -1.148953 | 0.139650 |
| 4 | 0.110501 | 0.692779 | 0.253224 |
| 5 | 1.260373 | 2.251253 | -1.638483 |

```
In [41]: df.iloc[2:, 1] = NA #Replac row 2 onwards, in col '1', with NA
df.iloc[4:, 2] = NA #Replac rows 4 onwards in col '2' with NA
df
```

Out[41]:

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 0 | -0.639978 | -0.799902 | -0.070594 |
| 1 | 0.725652 | 0.243492 | 0.511909 |
| 2 | -0.386548 | NaN | 0.272159 |
| 3 | -0.475659 | NaN | 0.139650 |
| 4 | 0.110501 | NaN | NaN |
| 5 | 1.260373 | NaN | NaN |


```
In [42]: # can also do
df.fillna(method = 'ffill')
```

```
Out[42]:
```

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 0 | -0.639978 | -0.799902 | -0.070594 |
| 1 | 0.725652 | 0.243492 | 0.511909 |
| 2 | -0.386548 | 0.243492 | 0.272159 |
| 3 | -0.475659 | 0.243492 | 0.139650 |
| 4 | 0.110501 | 0.243492 | 0.139650 |
| 5 | 1.260373 | 0.243492 | 0.139650 |

ffill() function is used to fill the missing value in the dataframe. 'ffill' stands for 'forward fill' and will propagate last valid observation forward

AKA repeats the NA with the previous data value

```
In [43]: # Only apply the Forward Fill to 2 values, max
df.fillna(method = 'ffill', limit = 2)
```

```
Out[43]:
```

| | 0 | 1 | 2 |
|---|-----------|-----------|-----------|
| 0 | -0.639978 | -0.799902 | -0.070594 |
| 1 | 0.725652 | 0.243492 | 0.511909 |
| 2 | -0.386548 | 0.243492 | 0.272159 |
| 3 | -0.475659 | 0.243492 | 0.139650 |
| 4 | 0.110501 | NaN | 0.139650 |
| 5 | 1.260373 | NaN | 0.139650 |

```
In [44]: # You can also pass the mean or median value with ffill to a Series!
data = pd.Series([1., NA, 3.5, NA, 7])
data # what the series looks like before:
```

```
Out[44]: 0    1.0
1    NaN
2    3.5
3    NaN
4    7.0
dtype: float64
```

```
In [45]: #What it Looks Like after, when we replace all NAs with the median value  
data.fillna(data.mean())
```

```
Out[45]: 0    1.000000  
        1    3.833333  
        2    3.500000  
        3    3.833333  
        4    7.000000  
        dtype: float64
```

See Table 7.2 on pg 197 for more fillna arguments

7.2 Data Transformation

Removing Duplicates

```
In [46]: # Example of a df with duplicates:  
data = pd.DataFrame({'k1': ['one', 'two'] * 3 + ['two'],  
                    'k2': [1, 1, 2, 3, 3, 4, 4]})  
data
```

```
Out[46]:
```

| | k1 | k2 |
|---|-----|----|
| 0 | one | 1 |
| 1 | two | 1 |
| 2 | one | 2 |
| 3 | two | 3 |
| 4 | one | 3 |
| 5 | two | 4 |
| 6 | two | 4 |

```
In [47]: # The df method 'duplicated' returns a boolean Series indication if there are
         any duplicates in the NAME of the row, based on the PREVIOUS one... i.e. so i
         f there are 2 in a row
         data.duplicated()
```

```
Out[47]: 0    False
         1    False
         2    False
         3    False
         4    False
         5    False
         6     True
         dtype: bool
```

```
In [48]: # drop_duplicates returns a df where the 'duplicated' array is False:
         data.drop_duplicates()

         #aka. this method will drop all cases where method 'duplicated' is True.
```

```
Out[48]:
```

| | k1 | k2 |
|---|-----|----|
| 0 | one | 1 |
| 1 | two | 1 |
| 2 | one | 2 |
| 3 | two | 3 |
| 4 | one | 3 |
| 5 | two | 4 |

You can also indicate if you only want to drop duplicates from a specific column!

```
In [49]: # First, start by adding another col to the df, called v1
         data['v1'] = range(7)
         data
```

```
Out[49]:
```

| | k1 | k2 | v1 |
|---|-----|----|----|
| 0 | one | 1 | 0 |
| 1 | two | 1 | 1 |
| 2 | one | 2 | 2 |
| 3 | two | 3 | 3 |
| 4 | one | 3 | 4 |
| 5 | two | 4 | 5 |
| 6 | two | 4 | 6 |

```
In [50]: # Next, drop any duplicates that are found in k1 only
data.drop_duplicates(['k1'])

# Remember that when deletings values like NAs or duplicates, this will drop t
he ENTIRE row it is found in, so that the data is neat/matches an even matrix/
dataframe form.
```

Out[50]:

| | k1 | k2 | v1 |
|---|-----|----|----|
| 0 | one | 1 | 0 |
| 1 | two | 1 | 1 |

Transforming data using a function or mapping

You may want to transform values in an array, Series, or column in DF.

```
In [6]: # First, create some data to work with

data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',
                              'Pastrami', 'corned beef', 'Bacon',
                              'pastrami', 'honey ham', 'nova lox'],
                    'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})

data
```

Out[6]:

| | food | ounces |
|---|-------------|--------|
| 0 | bacon | 4.0 |
| 1 | pulled pork | 3.0 |
| 2 | bacon | 12.0 |
| 3 | Pastrami | 6.0 |
| 4 | corned beef | 7.5 |
| 5 | Bacon | 8.0 |
| 6 | pastrami | 3.0 |
| 7 | honey ham | 5.0 |
| 8 | nova lox | 6.0 |

```
In [4]: # If you want to add a column that tells the type of animal of each good.
# First, we can write down the _mapping_ or legend for these meatz

meat_to_animal = {
    'bacon': 'pig',
    'pulled pork': 'pig',
    'pastrami': 'cow',
    'corned beef': 'cow',
    'honey ham': 'pig',
    'nova lox': 'salmon'
}
```

Mapping!

Python's `map()` is a built-in function that allows you to process and transform all the items in an iterable without using an explicit for loop, a technique commonly known as mapping. `map()` is useful when you need to apply a transformation function to each item in an iterable and transform them into a new iterable

<https://realpython.com/python-map-function/>

In other words, the map is like adding a "label" to X variables/categories in that label to allow for transformations to be applied to just that category. See example below.

```
In [7]: # The 'map' method on a Series accepts a function or dict-like object containing a map...
# BUT before that, the capitalisations needs to match EXACTLY.
# So, we need to transform the data so capitalisations match.

lowercased = data['food'].str.lower()
lowercased
```

```
Out[7]: 0      bacon
1  pulled pork
2      bacon
3    pastrami
4  corned beef
5      bacon
6    pastrami
7  honey ham
8    nova lox
Name: food, dtype: object
```

```
In [55]: # Now, we are going to APPLY this new Lower case key to the data + the map...

data['animal'] = lowercased.map(meat_to_animal) # Create the key/column 'animal', which will be the lowercased MAP values of meat_to_animal
data

# See page 199 for a better, color-coated version of the table. There you will see that animal col looks different since it is a _map_
```

Out[55]:

| | food | ounces | animal |
|---|-------------|--------|--------|
| 0 | bacon | 4.0 | pig |
| 1 | pulled pork | 3.0 | pig |
| 2 | bacon | 12.0 | pig |
| 3 | Pastrami | 6.0 | cow |
| 4 | corned beef | 7.5 | cow |
| 5 | Bacon | 8.0 | pig |
| 6 | pastrami | 3.0 | cow |
| 7 | honey ham | 5.0 | pig |
| 8 | nova lox | 6.0 | salmon |

[!] WTF does this mean?? pg. 300

fillna returns a `_new object_`, but you can modify the existing object in place...
(??)

<https://realpython.com/python-lambda/> (<https://realpython.com/python-lambda/>)

```
In [56]: # Instead of the above, a function could have also been passed instead

data['food'].map(lambda x: meat_to_animal[x.lower()])

#the above is saying...???? BUT WHY x. lower() ?! pg 200
# Answer: for the 'food' col in data df, add the map to x, of the 'meat_to_animal',
```

```
Out[56]: 0      pig
         1      pig
         2      pig
         3      cow
         4      cow
         5      pig
         6      cow
         7      pig
         8  salmon
         Name: food, dtype: object
```

Replacing Values

```
In [8]: # First, create example data/Series
data = pd.Series([1., -999., 2., -999., -1000., 3.])
data

#In this example -999 could stand for missing data.
```

```
Out[8]: 0      1.0
         1    -999.0
         2      2.0
         3    -999.0
         4   -1000.0
         5      3.0
         dtype: float64
```

```
In [9]: # Now, we want to replace -999 values with NA/NaN
data.replace(-999, np.nan)
```

```
Out[9]: 0      1.0
         1     NaN
         2      2.0
         3     NaN
         4   -1000.0
         5      3.0
         dtype: float64
```

```
In [10]: # IF you want to replace MULTIPLE values at the same time, pass a list!
data.replace([-999, -1000], np.nan)
```

```
Out[10]: 0    1.0
         1    NaN
         2    2.0
         3    NaN
         4    NaN
         5    3.0
         dtype: float64
```

```
In [11]: # to place MULTIPLE values with THEIR OWN unique value:
data.replace([-999, -1000], [np.nan, 0]) # Where np.nan is for -999, and 0 for -1000
```

```
Out[11]: 0    1.0
         1    NaN
         2    2.0
         3    NaN
         4    0.0
         5    3.0
         dtype: float64
```

```
In [12]: # The above can ALSO be passed as a dict
data.replace({-999: np.nan, -1000: 0})
```

```
Out[12]: 0    1.0
         1    NaN
         2    2.0
         3    NaN
         4    0.0
         5    3.0
         dtype: float64
```

Renaming Axis Indexes

Axis labels can also be transformed with a function or mapping. This will produce a new, differently labeled object. You can also modify axes in-place and without creating a new data structure.

```
In [2]: data = pd.DataFrame(np.arange(12).reshape((3, 4)),
                           index = ['Ohio', 'Colorado', 'New York'],
                           columns = ['one', 'two', 'three', 'four'])
data
```

```
Out[2]:
```

| | one | two | three | four |
|----------|-----|-----|-------|------|
| Ohio | 0 | 1 | 2 | 3 |
| Colorado | 4 | 5 | 6 | 7 |
| New York | 8 | 9 | 10 | 11 |


```
In [3]: # Axis indexes also have a map method:
transform = lambda x: x[:4].upper()

# First, we are creating the function, which states:
''' The function called transform does this:
takes the input value(x), up to the first 4 observations of the
upper half of the input value (x)'''
```

```
Out[3]: ' The function called transform does this:\ntakes the input value(x), up to t
he first 4 observations of the \nupper half of the input value (x)'
```

```
In [4]: #Now, apply the lambda function 'transform'
data.index = data.index.map(transform)
```

```
In [5]: # Now Look at the data
data

# Here, the 'transform' function was successful, and replaced the previous axi
s names with shorter ones.
```

Out[5]:

| | one | two | three | four |
|------|-----|-----|-------|------|
| OHIO | 0 | 1 | 2 | 3 |
| COLO | 4 | 5 | 6 | 7 |
| NEW | 8 | 9 | 10 | 11 |

```
In [6]: # Using the method 'rename' will transform the data, but NOT modify the origin
al data
data.rename(index = str.title, columns = str.upper)

# The above just changes the title of the index to uppercase
```

Out[6]:

| | ONE | TWO | THREE | FOUR |
|------|-----|-----|-------|------|
| Ohio | 0 | 1 | 2 | 3 |
| Colo | 4 | 5 | 6 | 7 |
| New | 8 | 9 | 10 | 11 |

```
In [7]: # Rename can also be like with a dict-like object, and new axis labels
data.rename(index = {'OHIO': 'INDIANA'},
            columns = {'three': 'peekaboo'})

# The above renames the index Ohio to Indian, and col three to peekaboo
```

Out[7]:

| | one | two | peekaboo | four |
|---------|-----|-----|----------|------|
| INDIANA | 0 | 1 | 2 | 3 |
| COLO | 4 | 5 | 6 | 7 |
| NEW | 8 | 9 | 10 | 11 |

Rename is used when you DO NOT want to copy the Dataframe manually and assign values to its index and cols.

If you want to modify the data with rename officially, you can pass 'inplace = True'.

```
In [8]: data.rename(index = {'OHIO': 'INDIANA'}, inplace = True)
data
```

Out[8]:

| | one | two | three | four |
|---------|-----|-----|-------|------|
| INDIANA | 0 | 1 | 2 | 3 |
| COLO | 4 | 5 | 6 | 7 |
| NEW | 8 | 9 | 10 | 11 |

Discretization and Binning

Continuous data is usually 'discretized' aka, separated into bins. See example below with groups of people in a study and you want to group them into discrete age buckets:

```
In [12]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
bins = [18, 25, 35, 60, 100]

# Now divide the ages data into bins, using pd.cut

cats = pd.cut(ages, bins) # This says put the data along with its bin value.
cats

# What is returned is a special pandas 'Categorical' objects, and can be treated like an array of strings.

Out[12]: [(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100],
(35, 60], (35, 60], (25, 35]]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

See official doc on pd.cut aka, using bins in pandas: <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.cut.html> (<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.cut.html>)

An example project with pd.cut / bins:

<https://realpython.com/fast-flexible-pandas/> (<https://realpython.com/fast-flexible-pandas/>)

```
In [13]: cats.codes # which bin each data from 'ages' belongs to...the first, second, etc.
```

```
Out[13]: array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)
```

```
In [14]: cats.categories # all the different bins that are observed for the cats/ages data
```

```
Out[14]: IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]],
                        closed='right',
                        dtype='interval[int64]')
```

```
In [17]: pd.value_counts(cats) # the bin counts for the results of pd.cuts, like a bin tally
```

```
Out[17]: (18, 25]      5
         (35, 60]      3
         (25, 35]      3
         (60, 100]     1
         dtype: int64
```

Consistent with mathematical notation for intervals, a parenthesis means that the side is open, while the square bracket means it is closed (inclusive). You can change which side is closed by passing `right=False`: pg. 203

```
In [19]: pd.cut(ages, [18, 26, 36, 61, 100], right = False) # see above for this to make sense, now it is reverse which bin is open, which is closed.
```

```
Out[19]: [[18, 26), [18, 26), [18, 26), [26, 36), [18, 26), ..., [26, 36), [61, 100),
         [36, 61), [36, 61), [26, 36)]
         Length: 12
         Categories (4, interval[int64]): [[18, 26) < [26, 36) < [36, 61) < [61, 100)]
```

You can also pass your OWN bin names, by passing a list or array to the 'labels' option!

```
In [20]: group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'Senior']
         pd.cut(ages, bins, labels = group_names)
```

```
Out[20]: [Youth, Youth, Youth, YoungAdult, Youth, ..., YoungAdult, Senior, MiddleAged,
         MiddleAged, YoungAdult]
         Length: 12
         Categories (4, object): [Youth < YoungAdult < MiddleAged < Senior]
```

If you pass an integer number of bins to cut instead of explicit bin edges, it will compute equal-length bins based on the minimum and maximum values in the data. Consider the case of some uniformly distributed data chopped into fourths: (see below) pg. 204

In [22]: `data = np.random.rand(20)`
`data`

Out[22]: `array([0.79823906, 0.75341494, 0.09454416, 0.39304138, 0.10175479,`
`0.95731211, 0.57518799, 0.38454512, 0.5360092 , 0.98710699,`
`0.43339816, 0.70082235, 0.36679895, 0.40754292, 0.35509414,`
`0.64113761, 0.75310718, 0.74036004, 0.81134476, 0.53548824])`

In [24]: `pd.cut(data, 4, precision = 2) # precision limits the decimal place to 2`

Out[24]: `[(0.76, 0.99], (0.54, 0.76], (0.094, 0.32], (0.32, 0.54], (0.094, 0.32], ...,`
`(0.54, 0.76], (0.54, 0.76], (0.54, 0.76], (0.76, 0.99], (0.32, 0.54]]`
`Length: 20`
`Categories (4, interval[float64]): [(0.094, 0.32] < (0.32, 0.54] < (0.54, 0.7`
`6] < (0.76, 0.99]]`

[?] Does this mean that you should NOT use integer number of bins?

Using qcut

There is also the function 'qcut' which will bin the data based on sample quantiles. You may want to use this instead of 'cut' because 'cut' may not give you an even number of data points in each bin.

Since qcut uses sample quantiles instead, by definition you will obtain roughly equal-size bins: (see below)

```
In [25]: data = np.random.randn(1000) # Normally distributed
cats = pd.qcut(data, 4) # Cut into 4 quartiles, approx. evenly.

cats
```

```
Out[25]: [(0.0267, 0.695], (-0.629, 0.0267], (-0.629, 0.0267], (-3.046, -0.629], (0.69
5, 2.716], ..., (-0.629, 0.0267], (-3.046, -0.629], (-0.629, 0.0267], (-0.62
9, 0.0267], (-3.046, -0.629]]
Length: 1000
Categories (4, interval[float64]): [(-3.046, -0.629] < (-0.629, 0.0267] < (0.
0267, 0.695] < (0.695, 2.716]]
```

```
In [26]: pd.value_counts(cats) # Shows you how many data points are in each 'bin'
```

```
Out[26]: (0.695, 2.716]      250
(0.0267, 0.695]      250
(-0.629, 0.0267]      250
(-3.046, -0.629]      250
dtype: int64
```

```
In [29]: # You can also pass your OWN quantiles with qcut, using no. between 0 and 1, i
nclusive.
dog = pd.qcut(data, [0, 0.1, 0.5, 0.9, 1.])
dog
```

```
Out[29]: [(0.0267, 1.242], (-1.3, 0.0267], (-1.3, 0.0267], (-1.3, 0.0267], (0.0267, 1.
242], ..., (-1.3, 0.0267], (-3.046, -1.3], (-1.3, 0.0267], (-1.3, 0.0267], (-
1.3, 0.0267]]
Length: 1000
Categories (4, interval[float64]): [(-3.046, -1.3] < (-1.3, 0.0267] < (0.026
7, 1.242] < (1.242, 2.716]]
```

```
In [30]: pd.value_counts(dog) # this is how it looks...
```

```
Out[30]: (0.0267, 1.242]      400
(-1.3, 0.0267]      400
(1.242, 2.716]      100
(-3.046, -1.3]      100
dtype: int64
```

[?] When and why would you ever want to set your own quantiles? I dont get the example either... (see above chunks).

Detecting and Filtering Outliers

Mailly just applying array operations. Below will be an example with normally distributed data.

```
In [31]: data = pd.DataFrame(np.random.randn(1000, 4))
data.describe()
```

```
Out[31]:
```

| | 0 | 1 | 2 | 3 |
|-------|-------------|-------------|-------------|-------------|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 |
| mean | 0.045703 | -0.006905 | -0.035234 | -0.031098 |
| std | 1.003481 | 0.999982 | 0.975964 | 1.001750 |
| min | -3.004345 | -3.727977 | -3.728292 | -3.228248 |
| 25% | -0.615337 | -0.684660 | -0.667265 | -0.687612 |
| 50% | 0.058913 | 0.044233 | -0.006272 | -0.035844 |
| 75% | 0.672576 | 0.668767 | 0.597902 | 0.659980 |
| max | 3.164499 | 2.834683 | 2.631226 | 2.802174 |

```
In [32]: #Suppose you wanted to find values in one of the columns exceeding 3 in absolute value:
col = data[2] # first isolate/ select only the 3rd column of the data

col[np.abs(col) > 3]
```

```
Out[32]: 599    -3.728292
Name: 2, dtype: float64
```

```
In [33]: # To select the rows that have a value of 3 or -3 you can use the 'any' method...
data[(np.abs(data) > 3).any(1)]

# Notes, SELECT A ROW, THE WHOLE DAMN ROW that has a +-3 value
```

```
Out[33]:
```

| | 0 | 1 | 2 | 3 |
|-----|-----------|-----------|-----------|-----------|
| 143 | -3.004345 | -0.899828 | 0.240154 | 0.192567 |
| 400 | 3.164499 | 0.813197 | -0.078097 | -0.080511 |
| 426 | 3.002462 | -1.296210 | -0.842690 | 1.191065 |
| 503 | -1.410977 | -3.727977 | 0.805915 | 0.865970 |
| 599 | -1.009860 | -0.396570 | -3.728292 | -1.332789 |
| 626 | 2.153357 | -1.127845 | 0.182386 | -3.228248 |

```
In [37]: # 'Code to cap values outside the interval -3 to 3'
data[np.abs(data) > 3] = np.sign(data) * 3 # we do this to get just a clean -
3 or 3
data.describe()

## ??? But what is the point of doing this
```

Out[37]:

| | 0 | 1 | 2 | 3 |
|-------|-------------|-------------|-------------|-------------|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 |
| mean | 0.045541 | -0.006177 | -0.034506 | -0.030870 |
| std | 1.002962 | 0.997532 | 0.973474 | 1.001047 |
| min | -3.000000 | -3.000000 | -3.000000 | -3.000000 |
| 25% | -0.615337 | -0.684660 | -0.667265 | -0.687612 |
| 50% | 0.058913 | 0.044233 | -0.006272 | -0.035844 |
| 75% | 0.672576 | 0.668767 | 0.597902 | 0.659980 |
| max | 3.000000 | 2.834683 | 2.631226 | 2.802174 |

```
In [38]: # np.sign(data) produced 1 and -1 values on whether the values in the data are
pos or negative
np.sign(data).head()
```

Out[38]:

| | 0 | 1 | 2 | 3 |
|---|------|-----|------|------|
| 0 | -1.0 | 1.0 | -1.0 | -1.0 |
| 1 | -1.0 | 1.0 | 1.0 | -1.0 |
| 2 | 1.0 | 1.0 | 1.0 | 1.0 |
| 3 | 1.0 | 1.0 | 1.0 | 1.0 |
| 4 | -1.0 | 1.0 | -1.0 | -1.0 |

Permutation and Random Sampling

Permuting (randomly reordering) can be done with `np.random.permutation`. Calling permutation with the length of the axis you want to permute creates an array of integers indicating the new reordering: (see below).

Also for offic. doc.

<https://numpy.org/doc/stable/reference/random/generated/numpy.random.permutation.html>

(<https://numpy.org/doc/stable/reference/random/generated/numpy.random.permutation.html>)

```
In [39]: df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))
df

#Creating the data and taking a preliminary look at the data
```

```
Out[39]:
```

| | 0 | 1 | 2 | 3 |
|---|----|----|----|----|
| 0 | 0 | 1 | 2 | 3 |
| 1 | 4 | 5 | 6 | 7 |
| 2 | 8 | 9 | 10 | 11 |
| 3 | 12 | 13 | 14 | 15 |
| 4 | 16 | 17 | 18 | 19 |

```
In [40]: sampler = np.random.permutation(5) # Randomly rearrange an array of 5 items
sampler # This outputs says re-arrange items 0, 1, 2, 3, 4 like so:
```

```
Out[40]: array([2, 0, 3, 4, 1])
```

```
In [41]: # The array can then be used in _iloc_ based indexing or the equivalent _take_
df
```

```
Out[41]:
```

| | 0 | 1 | 2 | 3 |
|---|----|----|----|----|
| 0 | 0 | 1 | 2 | 3 |
| 1 | 4 | 5 | 6 | 7 |
| 2 | 8 | 9 | 10 | 11 |
| 3 | 12 | 13 | 14 | 15 |
| 4 | 16 | 17 | 18 | 19 |

```
In [42]: df.take(sampler) # This randomly permuted the INDEXES aka the ENTIRE ROWS of
the data
```

```
Out[42]:
```

| | 0 | 1 | 2 | 3 |
|---|----|----|----|----|
| 2 | 8 | 9 | 10 | 11 |
| 0 | 0 | 1 | 2 | 3 |
| 3 | 12 | 13 | 14 | 15 |
| 4 | 16 | 17 | 18 | 19 |
| 1 | 4 | 5 | 6 | 7 |

[?] So, does that mean sampler MUST match the no. of rows/indexes in the dataframe to work?

In [43]: *# To select a random subset w.o. replacement, you can use _sample_ method on a Series or DF*
`df.sample(n = 3) # Give me 3 random rows/indexes`

Out[43]:

| | 0 | 1 | 2 | 3 |
|---|---|---|----|----|
| 1 | 4 | 5 | 6 | 7 |
| 0 | 0 | 1 | 2 | 3 |
| 2 | 8 | 9 | 10 | 11 |

In [45]: *# To generate a sample WITH replacement, pass replace = True to sample.*
`choices = pd.Series([5, 7, -1, 6, 4])`
`draws = choices.sample(n = 10, replace = True)`
`draws`

Out[45]:

| | |
|---|----|
| 0 | 5 |
| 4 | 4 |
| 3 | 6 |
| 2 | -1 |
| 3 | 6 |
| 1 | 7 |
| 2 | -1 |
| 2 | -1 |
| 0 | 5 |
| 0 | 5 |

dtype: int64

Computing Indicator/Dummy Variables

Pandas has a `get_dummies` function, and converts categorical variables into dummy or indicator variables.

In [47]: `df = pd.DataFrame({'key': ['b', 'b', 'a', 'c', 'a', 'b'],`
`'data1': range(6)})`

`df`

Out[47]:

| | key | data1 |
|---|-----|-------|
| 0 | b | 0 |
| 1 | b | 1 |
| 2 | a | 2 |
| 3 | c | 3 |
| 4 | a | 4 |
| 5 | b | 5 |

In [48]: `pd.get_dummies(df['key'])` # Because there are 3 diff values, a, b, c each have their own key...
 # Where col a means give a 0 to all values that are NOT a, and col b says give value 0 to all those thare NOT b, etc.

Out[48]:

| | a | b | c |
|---|---|---|---|
| 0 | 0 | 1 | 0 |
| 1 | 0 | 1 | 0 |
| 2 | 1 | 0 | 0 |
| 3 | 0 | 0 | 1 |
| 4 | 1 | 0 | 0 |
| 5 | 0 | 1 | 0 |

In [50]: # To make it easier to read, can add a prefix to name the different keys better.
`dummies = pd.get_dummies(df['key'], prefix = 'key')`
`df_with_dummy = df[['data1']].join(dummies)`
`df_with_dummy`

Out[50]:

| | data1 | key_a | key_b | key_c |
|---|-------|-------|-------|-------|
| 0 | 0 | 0 | 1 | 0 |
| 1 | 1 | 0 | 1 | 0 |
| 2 | 2 | 1 | 0 | 0 |
| 3 | 3 | 0 | 0 | 1 |
| 4 | 4 | 1 | 0 | 0 |
| 5 | 5 | 0 | 1 | 0 |

If a row in a DF belongs to multipl categories, it is more complicated. See example below looking at the MovieLens 1M dataset, which is investigated in more detail in Ch. 14

```
In [2]: mnames = ['movie_id', 'title', 'genres']

movies = pd.read_table('C:\\Users\\Kitty\\Desktop\\learnpy\\movies.dat', sep =
'::',
                        header = None, names = mnames) # Had to put r in fron
t, or \\

movies[:10] # give me the first 10 rows of data
```

<ipython-input-2-0c96f39157f3>:3: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
movies = pd.read_table('C:\\Users\\Kitty\\Desktop\\learnpy\\movies.dat', se
p = '::',
```

Out[2]:

| | movie_id | title | genres |
|---|----------|------------------------------------|------------------------------|
| 0 | 1 | Toy Story (1995) | Animation Children's Comedy |
| 1 | 2 | Jumanji (1995) | Adventure Children's Fantasy |
| 2 | 3 | Grumpier Old Men (1995) | Comedy Romance |
| 3 | 4 | Waiting to Exhale (1995) | Comedy Drama |
| 4 | 5 | Father of the Bride Part II (1995) | Comedy |
| 5 | 6 | Heat (1995) | Action Crime Thriller |
| 6 | 7 | Sabrina (1995) | Comedy Romance |
| 7 | 8 | Tom and Huck (1995) | Adventure Children's |
| 8 | 9 | Sudden Death (1995) | Action |
| 9 | 10 | GoldenEye (1995) | Action Adventure Thriller |

```
In [16]: # To add indicator variables, need to wrangle more.
#First, extract the list of unique genres in the dataset

all_genres = []

for x in movies.genres:
    all_genres.extend(x.split('|'))

# Above- creating a for loop to go through all genres and split them up by the
```

```
In [17]: genres = pd.unique(all_genres) #print all the unique cases of genres, to see h
ow many different ones there are

genres
```

```
Out[17]: array(['Animation', 'Children's', 'Comedy', 'Adventure', 'Fantasy',
'Romance', 'Drama', 'Action', 'Crime', 'Thriller', 'Horror',
'Sci-Fi', 'Documentary', 'War', 'Musical', 'Mystery', 'Film-Noir',
'Western'], dtype=object)
```

One way to construct the indicator DF is to start with a DF of all zeros (see code below)

[?] Is there another way to do this?

```
In [19]: zero_matrix = np.zeros((len(movies), len(genres))) # make a 0 matrix that has
          same length/dimensions as the following variables in the DF
          dummies = pd.DataFrame(zero_matrix, columns = genres)

          dummies
```

Out[19]:

| | Animation | Children's | Comedy | Adventure | Fantasy | Romance | Drama | Action | Crime | Th |
|------|-----------|------------|--------|-----------|---------|---------|-------|--------|-------|-----|
| 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 1 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 2 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 3 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 4 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 3878 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 3879 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 3880 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 3881 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |
| 3882 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |

3883 rows × 18 columns

Now interact through each movie and set in each row of dummies to 1... by using dummies.columns to calculate the column indices for each genre.

```
In [24]: gen = movies.genres[0] # the genre of the movies in the first row
          gen

          # This is kind of like setting the stage/starting pt for creating the dummie v
          ariables
```

Out[24]: "Animation|Children's|Comedy"

```
In [25]: gen.split('|') #split the gen variable
```

Out[25]: ['Animation', 'Children's', 'Comedy']

[?] WTF does this do, pg 210. pls explain

```
In [4]: dummies.columns.get_indexer(gen.split('|'))
```

```
# WTF HAPPENED BRUH
```

```
-----
NameError                                Traceback (most recent call last)
```

```
<ipython-input-4-ae17c8654c4a> in <module>
```

```
----> 1 dummies.columns.get_indexer(gen.split('|'))
```

```
2
```

```
3 # WTF HAPPENED BRUH
```

```
NameError: name 'dummies' is not defined
```

[?] WTF happened here pg 209. pls explain

```
In [3]: # Next, can use .iloc to set values based on these indicies that were jsut cre
ated
for i, gen in enumerate(movies.genres):
    indicies = dummies.columns.get_indexer(gen.split('|'))
    dummies.iloc[i, indicies] = 1
```

```
-----
NameError                                Traceback (most recent call last)
```

```
<ipython-input-3-572db78995b7> in <module>
```

```
1 # Next, can use .iloc to set values based on these indicies that were
jsut created
```

```
2 for i, gen in enumerate(movies.genres):
```

```
----> 3     indicies = dummies.columns.get_indexer(gen.split('|'))
```

```
4     dummies.iloc[i, indicies] = 1
```

```
NameError: name 'dummies' is not defined
```

```
In [31]: # Now combine with movies...
movies_windic = movies.join(dummies.add_prefix('Genre_'))
movies_windic.iloc[0]
```

```
Out[31]: movie_id          1
title          Toy Story (1995)
genres        Animation|Children's|Comedy
Genre_Animation          1
Genre_Children's          1
Genre_Comedy              1
Genre_Adventure           0
Genre_Fantasy             0
Genre_Romance             0
Genre_Drama               0
Genre_Action              0
Genre_Crime               0
Genre_Thriller            0
Genre_Horror              0
Genre_Sci-Fi             0
Genre_Documentary        0
Genre_War                 0
Genre_Musical             0
Genre_Mystery             0
Genre_Film-Noir           0
Genre_Western             0
Name: 0, dtype: object
```

Another useful method for stats applications is to combine `get_dummies` with a discretization function like 'cut':

```
In [32]: np.random.seed(12345)
values = np.random.rand(10)
values
```

```
Out[32]: array([0.92961609, 0.31637555, 0.18391881, 0.20456028, 0.56772503,
                0.5955447 , 0.96451452, 0.6531771 , 0.74890664, 0.65356987])
```

```
In [34]: bins = [0, 0.2, 0.4, 0.6, 0.8, 1]
pd.get_dummies(pd.cut(values, bins))

#Set the random seed with numpy.random.seed to make the example 'deterministic'
#Later, the book will explore pd.get_dummies
```

Out[34]:

| | (0.0, 0.2] | (0.2, 0.4] | (0.4, 0.6] | (0.6, 0.8] | (0.8, 1.0] |
|---|------------|------------|------------|------------|------------|
| 0 | 0 | 0 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 | 0 | 0 |
| 2 | 1 | 0 | 0 | 0 | 0 |
| 3 | 0 | 1 | 0 | 0 | 0 |
| 4 | 0 | 0 | 1 | 0 | 0 |
| 5 | 0 | 0 | 1 | 0 | 0 |
| 6 | 0 | 0 | 0 | 0 | 1 |
| 7 | 0 | 0 | 0 | 1 | 0 |
| 8 | 0 | 0 | 0 | 1 | 0 |
| 9 | 0 | 0 | 0 | 1 | 0 |

7.3 String Manipulation

String Object Methods

```
In [5]: # E.g., a common-separate string can be broken into pieces with 'split'
val = 'a,b, guido'
val.split(',')
```

Out[5]: ['a', 'b', ' guido']

```
In [6]: # 'split' is often combined with 'strip' to trim white space, including line breaks
pieces = [x.strip() for x in val.split(',')]
pieces
```

Out[6]: ['a', 'b', 'guido']

```
In [7]: # These substrings could be added together with 2 colon delimiter using additi
on
# WHUT

first, second, third = pieces
first + '::' + second + '::' + third

# ... Interesting...but why? -See below, this is not practical or very 'Pytho
n' like.
```

```
Out[7]: 'a::b::guido'
```

```
In [8]: # However, the above is not practical...
# Faster to pass a list or tuple to the 'join' method on the string '::'
'::'.join(pieces)
```

```
Out[8]: 'a::b::guido'
```

```
In [9]: # Other methods deal with locating substrings
'guido' in val
```

```
Out[9]: True
```

```
In [41]: val.index(',')
```

```
Out[41]: 1
```

```
In [10]: val.find(':') # Returns -1 if the string isn't found
```

```
Out[10]: -1
```

```
In [11]: val.index(':') # Where as this just returns an error, compared to "val find"
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-11-2c016e7367ac> in <module>
----> 1 val.index(':')

ValueError: substring not found
```

```
In [12]: # 'Count' returns the no. of occurrences of a particular substring
val.count('')
```

```
Out[12]: 12
```

```
In [13]: # 'Replace' will subsistitute a value with the one u tell it to.
# Commonly used to delete patterns too, by just passing an empty string

val.replace(',', ' '::) # For every ',' replace it with '::'
```

```
Out[13]: 'a::b:: guido'
```



```
In [14]: val.replace(',', ' ') # For every ',' DELETE by replacing it with nothing
```

```
Out[14]: 'ab  guido'
```

See table 7.3 on pg 213 for a more Python built-in string methods

Regular Expressions

A way to search or match string patterns in a text. A single expression is called a 'regex', and is a string formed according to the the regular expression language.

Python has built in 're', and its functions fall into 3 categories: pattern matching, substitution, and splitting.

```
In [15]: # Ex. We want to split a string with a variable number of whitespace character
#         S,
#         in regex, whitespace characters a '\s+'
#         Whitespace includes: tabs, spaces, and newlines.

import re
```

```
In [16]: text = "foo      bar\t baz   \tqux"
text
```

```
Out[16]: 'foo      bar\t baz   \tqux'
```

```
In [18]: re.split('\s+', text) # Split up 'text', by the whitespace
```

```
Out[18]: ['foo', 'bar', 'baz', 'qux']
```

```
In [19]: # When re.split is called, the regex is first *compiled*, and then split is called.
#         You can compile a regex urself, to make a reusable regex object.

regex = re.compile('\s+')
regex.split(text)

# So basically, make a function of the compiling first (?)
## ?? What is the point? Faster? Asnwer: YES
```

```
Out[19]: ['foo', 'bar', 'baz', 'qux']
```

```
In [20]: # If you want to get a list of all patterns matching the regex, use findall.

regex.findall(text) # Show me a list of all the whitespace in 'text'
```

```
Out[20]: ['      ', '\t ', '   ', '\t']
```

'match' and 'search' are similar to 'findall'

'findall' returns ALL matches in a string.

'search' returns only the FIRST match.

'match' ONLY matches at the beginning of the string.

See below for an example with a block of text and a regex capable of identifying most email addresses.

```
In [21]: # JUST AN EXAMPLE GIVEN IN THE BOOK, pg 214
text = """Dave dave@google.com
Steve steve@gmail.com
Rob rob@gmail.com
Ryan ryan@yahoo.com
"""

pattern = r'[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,4}' #give me only email addresses, anything WITH an @ symbol

# re.IGNORECASE makes the regex case-insensitive
regex = re.compile(pattern, flags=re.IGNORECASE)
```

```
In [22]: regex.findall(text)

#So this is saying, use the previously define 'regex' that are an email addresses, and not jsut a name
```

```
Out[22]: ['dave@google.com', 'steve@gmail.com', 'rob@gmail.com', 'ryan@yahoo.com']
```

```
In [23]: # 'search' returns a special match object for the 1st email only.... huh
m = regex.search(text)
m

# And tells us where it is... >_> ????
```

```
Out[23]: <re.Match object; span=(5, 20), match='dave@google.com'>
```

```
In [24]: text[m.start():m.end()] # of the value of the search result (aka 'm') pls give me the full beginning and end of that string
```

```
Out[24]: 'dave@google.com'
```

```
In [25]: #regex.match returns None, as it onyl will match if the pattern occurs at the START of the string
print(regex.match(text))
```

None

```
In [26]: # Similarly, 'sub' will return a new string with occurrences of the pattern re
         # place by the string
         print(regex.sub('REDACTED', text))

         # AKA use sub to replace whatever meets the previously defined 'regex' with th
         # e indicated value/in this case REDACTED

Dave REDACTED
Steve REDACTED
Rob REDACTED
Ryan REDACTED
```

Suppose you wanted to find email addresses and simultaneously segment each address into its three components: username, domain name, and domain suffix. To do this, put parentheses around the parts of the pattern to segment:

pg. 215

```
In [27]: # Step 1 - COPIED FROM BOOK
         pattern = r'([A-Z0-9._%+-]+)@([A-Z0-9.-]+)\.([A-Z]{2,4})'

         # This is the defining first what we want in the example text above:
         # To segment the emails in 3 component.

In [28]: # Step 2
         regex = re.compile(pattern, flags = re.IGNORECASE)

         # We are compiling it now/defining the regex, and saying that upper or lower c
         # ase dont matter

In [29]: # Step 3 -
         # A 'match' object produced by this regex returns a tuple of the pattern compo
         # nents with method 'groups'
         m = regex.match('wesm@bright.net') # instead of our text list, we are applying
         # above steps to this new email address.
         m.groups()

         # We we can put it all together, to get what we originally wanted in the exam
         # ple

Out[29]: ('wesm', 'bright', 'net')

In [30]: # 'findall' returns of List of tuples when the pattern has groups.
         regex.findall(text) # now apply the above to our previous "text" with all oth
         # er emails

Out[30]: [('dave', 'google', 'com'),
          ('steve', 'gmail', 'com'),
          ('rob', 'gmail', 'com'),
          ('ryan', 'yahoo', 'com')]
```

```
In [31]: # 'sub' also has access to groups in each match using special symbols
# Like \1 (first matches group) or \2 (second matched group, etc)

print(regex.sub(r'Username: \1, Domain: \2, Suffix: \3', text))

# the above is saying, using the regex, add these labels to each of the matched groups.
# E.g., call the 1st group 'Username', etc.
```

```
Dave Username: dave, Domain: google, Suffix: com
Steve Username: steve, Domain: gmail, Suffix: com
Rob Username: rob, Domain: gmail, Suffix: com
Ryan Username: ryan, Domain: yahoo, Suffix: com
```

See table 7.3 pg. 216 for table of more regex methods

Vectorized String Functions in pandas

(e.g., for when having to clean data, and some strings have missing data)

```
In [33]: data = {'Dave': 'dave@google.com', 'Steve': 'steve@gmail.com',
                'Rob': 'rob@gmail.com', 'Wes': np.nan}
```

```
In [34]: data = pd.Series(data)
data
```

```
Out[34]: Dave      dave@google.com
Steve    steve@gmail.com
Rob      rob@gmail.com
Wes      NaN
dtype: object
```

```
In [35]: data.isnull() # Check if/where is there missing data
```

```
Out[35]: Dave      False
Steve    False
Rob      False
Wes      True
dtype: bool
```

You can apply string and regex methods can be applied (by passing a lambda or other function) to each value using 'data.map' but it will fail on the NA / null values.

To deal with this, Series has array-oriented methods for string operations that skip NA values! These are accessed via Series's 'str' attribute.

E.g., we can check whether each email address has 'gmail' in it with 'str.contains' (see below)

```
In [36]: data.str.contains('gmail')
```

```
Out[36]: Dave      False
         Steve     True
         Rob       True
         Wes       NaN
         dtype: object
```

```
In [37]: # Regex can be use too, along with any other 're' optionsl ike IGNORECASE
         pattern
```

```
Out[37]: '([A-Z0-9._%+-]+)@([A-Z0-9.-]+)\\.([A-Z]{2,4})'
```

```
In [38]: data.str.findall(pattern, flags = re.IGNORECASE)
```

```
Out[38]: Dave      [(dave, google, com)]
         Steve     [(steve, gmail, com)]
         Rob       [(rob, gmail, com)]
         Wes       NaN
         dtype: object
```

```
In [50]: # Are a few ways to do *vectorised element retrival*
         # Either with 'str.get' or index into the 'str' attribute
         matches = data.str.match(pattern, flags = re.IGNORECASE)
         matches
```

```
Out[50]: Dave      True
         Steve     True
         Rob       True
         Wes       NaN
         dtype: object
```

```
In [49]: # To access elements in the embedded lists, can pass an index to either of the
         matches.str.get(1)

         ## ** I GOT AN ERROR? but y
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-49-43eaae90677b> in <module>
      1 # To access elements in the embedded lists, can pass an index to either of these functions
----> 2 matches.str.get(1)
      3
      4 ## ** I GOT AN ERROR? but y

~\anaconda3\lib\site-packages\pandas\core\generic.py in __getattr__(self, name)
    5268         or name in self._accessors
    5269     ):
-> 5270         return object.__getattribute__(self, name)
    5271     else:
    5272         if self._info_axis._can_hold_identifiers_and_holds_name(name):

~\anaconda3\lib\site-packages\pandas\core\accessor.py in __get__(self, obj, cls)
    185         # we're accessing the attribute of the class, i.e., DataSeries
    186         return self._accessor
--> 187     accessor_obj = self._accessor(obj)
    188     # Replace the property with the accessor object. Inspired by:
    189     # http://www.pydanny.com/cached-property.html

~\anaconda3\lib\site-packages\pandas\core\strings.py in __init__(self, data)
    2039
    2040     def __init__(self, data):
-> 2041         self._inferred_dtype = self._validate(data)
    2042         self._is_categorical = is_categorical_dtype(data)
    2043         self._is_string = data.dtype.name == "string"

~\anaconda3\lib\site-packages\pandas\core\strings.py in _validate(data)
    2096
    2097         if inferred_dtype not in allowed_types:
-> 2098             raise AttributeError("Can only use .str accessor with string values!")
    2099         return inferred_dtype
    2100
```

AttributeError: Can only use .str accessor with string values!

```
In [45]: str(matches)
```

```
Out[45]: 'Dave      True\nSteve    True\nRob      True\nWes      NaN\ndtype: object'
```

```
In [ ]: matches.str[0]
```

```
In [47]: data.str[:5]
```

```
Out[47]: Dave      dave@  
Steve    steve  
Rob      rob@g  
Wes      NaN  
dtype: object
```

See table 7.5 for more pandas string methods

7.4 Conclusion

EFFECTIVE DATA PREP CAN SIGNIFICANTLY IMPROVE PRODUCTIVITY, ALLOWING FOR MORE TIME TO BE SPENT ON ANALYSING THE DATA AND LESS TIME GETTING IT READ FOR ANALYSIS.